

# PDTL: Parallel and Distributed Triangle Listing for Massive Graphs

Ilias Giechaskiel<sup>1</sup>   George Panagopoulos   Eiko Yoneki



<sup>1</sup>Currently at University of Oxford

September 3, 2015

## Motivation

- ▶ Graphs are becoming massive-scale
  - ▶ Billions of edges and vertices
- ▶ Triangle counting has important applications
  - ▶ Metrics for connectivity, density, quality of nodes

## Outcomes

- ▶ PDTL: Parallel and Distributed Triangle Listing for Massive Graphs
  - ▶ Provable Memory, Network, CPU, I/O bounds
  - ▶ Relies on and amends MGT (Hu et al., SIGMOD 2013)
  - ▶ Outperforms state-of-the-art by 2 – 4 $\times$  using fewer resources
  - ▶ Reasonable even compared to fast in-memory algorithms
- ▶ Bring external-memory considerations to distributed systems!

## Definition

Given a simple undirected graph  $G = (V, E)$ , list all  $\{v_1, v_2, v_3\}$  such that  $v_i \in V$  and  $(v_i, v_j) \in E$  exactly once

## Motivation

- ▶ Shortest non-trivial cycles and cliques
- ▶ Blackbox for density and connectivity metrics
  - ▶ Clustering coefficient, transitivity ratio
  - ▶ Triangular connectivity,  $k$ -truss
- ▶ Spam and fake account detection
- ▶ Massive graphs, so in-memory algorithms not sufficient

## Definition

Given  $G = (V, E)$ , define  $G^* = (V, E^*)$  with  $(u, v) \in E^*$  iff  $(u, v) \in E$  and either  $d(u) < d(v)$  or  $d(u) = d(v)$  with  $u < v$

## Motivation

- ▶ Better asymptotic runtime
- ▶ Identify  $\{u, v, w\}$  when  $u \prec v \prec w$  with  $(u, v, w)$ 
  - ▶  $u$  is the *cone vertex*
  - ▶  $(v, w)$  is the *pivot edge*




Figure: Oriented triangle

## Key Idea

1. Load next set of edges  $S$  of  $G^*$  into memory
  - ▶ Fill proportion of available memory  $M$
  - ▶ “All-or-nothing” requirement for each vertex
2. For all  $u \in V$ 
  - 2.1 Build hash structures on  $N(u)$
  - 2.2 Find all triangles with cone  $u$  and pivot in  $S$
  - 2.3 Release the hash structures

## Our Modifications

- ▶ Work on sorted arrays, not sets
- ▶ Ignore all-or-nothing
- ▶ Same complexity!

# Parallel and Distributed Triangle Listing (PDTL)

## Target Environment

$R$  machines,  $P$  processors/machine,  $M$  memory/processor

## Key Insights

- ▶ Graph partitioning requires random accesses
  - ▶ Or excessive network traffic
- ▶ Duplicate graph across *each* machine
  - ▶ But make I/O-efficient accesses
- ▶ Run (part of) MGT on each processor
  - ▶ Each processor responsible for different set of (contiguous) pivot edges
  - ▶ Small memory footprint

## Optimizations

- ▶ Orientation runs once and can be parallelized
- ▶ Load balance using number of in-edges

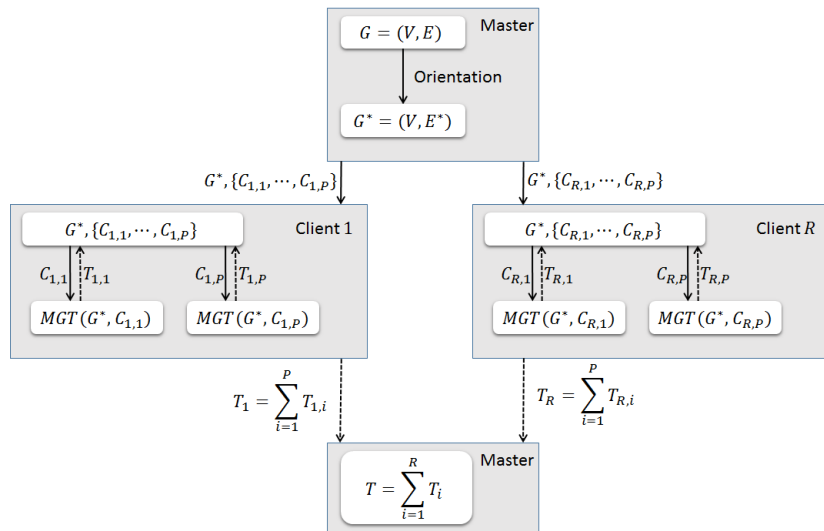


Figure: PDTL protocol

## Complexity

For triangle counting on all processors, with  $\alpha$  as the arboricity:

- ▶  $\Theta(RP + R|E|)$  Network traffic
- ▶  $\mathcal{O}\left(RP|E| + \frac{|E|^2}{M} + \alpha|E|\right)$  CPU computations
- ▶  $\mathcal{O}\left(RP\frac{|E|}{B} + \frac{|E|^2}{MB}\right)$  I/Os

## Key Decisions

- ▶ Orientation, load balancing parallelized and never re-computed
- ▶ Computation starts before transfers have finished
- ▶ Everything is included in analysis and timing



## Set-Up

- ▶ **Amazon EC2:** 4× c3.8xlarge, each with 32 CPUs, 60GB memory
- ▶ **Local Cluster:** 8 virtual nodes, each with 4 cores, 40GB memory
- ▶ **Local Multicore:** 24-core machine, 256GB memory

## Datasets

Graph	Nodes	Edges	Triangles	Size
soc-LiveJournal1	4.8M	68.0M	285,730,264	365MB
com-Orkut	3.1M	117.2M	627,584,181	917MB
Twitter	61.6M	1.5B	34,824,916,864	9.4GB
Yahoo	1.4B	6.6B	85,782,928,684	59GB
RMAT-26	67.1M	1.1B	51,559,452,522	8.4GB
RMAT-27	134.2M	2.1B	114,007,006,286	17GB
RMAT-28	268.4M	4.3B	251,913,686,661	34GB
RMAT-29	536.9M	8.6B	556,443,109,053	68GB

# Evaluation: PDTL Multicore Scalability





Figure: PDTL Total Time (without Orientation) in Local Multicore

# Evaluation: PDTL EC2 Scalability





Figure: PDTL Total Time (without Orientation) in Amazon EC2

# Evaluation: Load Balancing





Figure: PDTL Load Balancing

# Evaluation: Orientation

## Pre-processing time in seconds

Graphs	$d_{max}^*$	PDTL	PowerGraph	OPT
LiveJ1	687	1.4	-	106.8
Orkut	535	3.6	25.7	43.6
Twitter	4,102	32.8	233.2	437.6
Yahoo	1,540	235.6	-	-
RMAT-26	2,964	29.3	213.0	910.3




Figure: PDTL Multicore Orientation

## Distributed Frameworks

- ▶ PATRIC (Arifuzzaman et al., CIKM 2013) 564s, 200 cores, 4GB/core
  - ▶ PDTL faster using 8 cores and 1GB memory/core
  - ▶ PDTL 4× faster using 96 cores and 1GB memory/core
- ▶ CTTTP (Park et al., CIKM 2014) 5520s, 40 nodes, 4GB/node
  - ▶ PDTL with 1 core and 1GB memory (MGT): 2800s

## In-Memory Frameworks

- ▶ PDTL 141.8s using 128 cores and 128GB memory total
- ▶ (Sevenich et al., SNAKDD 2014)
  - ▶ 144.8s using 16\*2 Hyperthread Cores @ 2.2GHz and 264GB of memory
  - ▶ 91.5s using 128\*8 Hyperthread Cores @ 3.6GHz and 4TB of memory
- ▶ (Shun and Tangwongsan, ICDE 2015)
  - ▶ 55.9s using 40\*2 Hyperthread Cores @ 2.4GHz and 256GB of memory
  - ▶ 78.9s using 64 Cores @ 2.4 GHz and 188GB of memory

# Evaluation: PDTL More Comparisons

## PDTL and (Kim et al., SIGMOD 2014) performance in local multicore

Graph	PDTL		OPT		Speed-up	
	Orient (s)	Calc (s)	DB (s)	Calc (s)	Calc	Total
Twitter	32.8	262.9	235.2	437.6	1.66×	2.28×
RMAT-26	29.3	520.4	910.3	1011.2	1.94×	3.50×

## PDTL and (Gonzalez et al., OSDI 2012) performance in Amazon EC2

Graph	PDTL		PowerGraph		Speed-up	
	Calc (s)	Total (s)	Calc (s)	Total (s)	Calc	Total
Twitter	88.5	141.8	97.3	330.5	2.33×	1.86×
Yahoo	323.9	669.4	OOM	OOM	NA	NA
RMAT-26	138.6	180.6	176.7	389.7	2.16×	1.77×
RMAT-29	1533.5	1821.2	OOM	OOM	NA	NA

## Key Contributions

- ▶ Novel PDTL algorithm
  - ▶ Framework suitable for variety of environments
  - ▶ Provable CPU, I/O, Memory, and Network efficiency
- ▶ Improved state-of-the-art by  $2 - 4\times$  using fewer resources
  - ▶ Experimentally showed that algorithm is scalable
  - ▶ Orientation, balancing optimizations
  - ▶ Reasonable even compared to fast in-memory algorithms
- ▶ Focused on disk accesses in a distributed system

## Key Questions

- ▶ What other problems can benefit from external-memory, distributed algorithms?
- ▶ Your questions?
  - ▶ Reach me at [ilias.giechaskiel@cs.ox.ac.uk](mailto:ilias.giechaskiel@cs.ox.ac.uk)



## I/O Analysis

Let  $B$  denote the disk block size, and  $M$  the memory size

- ▶ Reading takes  $\text{scan}(N) = \Theta(N/B)$  I/Os
  - ▶ Orientation takes  $\text{scan}(|E|)$  I/Os and  $\mathcal{O}(|E|)$  CPU
- ▶ Sorting takes  $\text{sort}(N) = \mathcal{O}\left(N/B \log_{M/B} N/B\right)$  I/Os
  - ▶ Pre-sorting graph takes  $\text{sort}(|E|)$  I/Os and  $\mathcal{O}(|E| \log |E|)$  CPU

## Arboricity

The arboricity  $\alpha(G)$  of a graph  $G$  is the minimum number of edge-disjoint forests needed to cover the edges of  $G$ . It satisfies

$$\alpha(G) = \max_{G' \subseteq G} \left\lceil \frac{|E_{G'}|}{|V_{G'}| - 1} \right\rceil$$

where  $G'$  ranges over all subgraphs of  $G$  with  $\geq 2$  vertices



Shaikh Arifuzzaman, Maleq Khan, and Madhav Marathe. “PATRIC: A Parallel Algorithm for Counting Triangles in Massive Networks”. In: [CIKM 2013](#).



Deepayan Chakrabarti, Yiping Zhan, and Christos Faloutsos. “R-MAT: A recursive model for graph mining”. In: [SIAM 2004](#).



Joseph E. Gonzalez et al. “PowerGraph: Distributed Graph-parallel Computation on Natural Graphs”. In: [OSDI 2012](#).



Xiaocheng Hu, Yufei Tao, and Chin-Wan Chung. “Massive Graph Triangulation”. In: [SIGMOD 2013](#).



Jinha Kim et al. “OPT: A New Framework for Overlapped and Parallel Triangulation in Large-scale Graphs”. In: [SIGMOD 2014](#).



Haewoon Kwak et al. “What is Twitter, a social network or a news media?” In: WWW 2010.



Jure Leskovec. *SNAP: Stanford Large Network Dataset Collection*. <http://snap.stanford.edu/data/> 2009.



Ha-Myung Park et al. “MapReduce Triangle Enumeration With Guarantees”. In: CIKM 2014.



Martin Sevenich et al. “Fast In-Memory Triangle Listing for Large Real-World Graphs”. In: SNAKDD 2014.



J. Shun and K. Tangwongsan. “Multicore triangle computations without tuning”. In: ICDE 2015.



Yahoo! Webscope. *Yahoo! AltaVista Web Page Hyperlink Connectivity Graph, circa 2002*.  
<http://webscope.sandbox.yahoo.com/> 2009.



Zhi Yang et al. “Uncovering Social Network Sybils in the Wild”. In: KDD 2014.